

PufferFish: NUMA-Aware Work-stealing Library using Elastic Tasks

Vivek Kumar

IIIT New Delhi, India

Task Parallelism on Multicore Processors

```
1. int *A;
2. void Sort(int low, int high) {
3.     if((high-low)<LIMIT) return SeqSort(low, high);
4.     int Chunks=(high-low)/4;
5.
6.         Sort(/*Chunk C1*/);
7.         Sort(/*Chunk C2*/);
8.         Sort(/*Chunk C3*/);
9.         Sort(/*Chunk C4*/);
10.
11.
12.         Merge(/*Chunk C1*/, /*Chunk C2*/);
13.         Merge(/*Chunk C3*/, /*Chunk C4*/);
14.
15.     Merge(/*Chunk C12*/, /*Chunk C34*/);
16. }
```



Multicore Processor

Task Parallelism on Multicore Processors

```

1. int *A;
2. void Sort(int low, int high) {
3.   if((high-low)<LIMIT) return SeqSort(low, high);
4.   int Chunks=(high-low)/4;
5.   finish {
6.     async Sort(/*Chunk C1*/);
7.     async Sort(/*Chunk C2*/);
8.     async Sort(/*Chunk C3*/);
9.     async Sort(/*Chunk C4*/);
10.  }
11.  finish {
12.    async Merge(/*Chunk C1*/, /*Chunk C2*/);
13.    async Merge(/*Chunk C3*/, /*Chunk C4*/);
14.  }
15.  Merge(/*Chunk C12*/, /*Chunk C34*/);
16.}
    
```



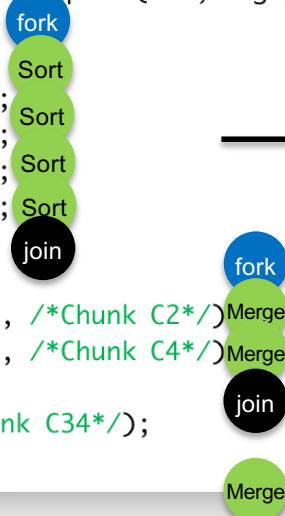
Multicore Processor

Serial elision
High productivity

Task Parallelism on Multicore Processors

```

1. int *A;
2. void Sort(int low, int high) {
3.   if((high-low)<LIMIT) return SeqSort(low, high);
4.   int Chunks=(high-low)/4;
5.   finish {
6.     async Sort(/*Chunk C1*/);
7.     async Sort(/*Chunk C2*/);
8.     async Sort(/*Chunk C3*/);
9.     async Sort(/*Chunk C4*/);
10.  }
11.  finish {
12.    async Merge(/*Chunk C1*/, /*Chunk C2*/);
13.    async Merge(/*Chunk C3*/, /*Chunk C4*/);
14.  }
15.  Merge(/*Chunk C12*/, /*Chunk C34*/);
16.}
    
```



Multicore Processor

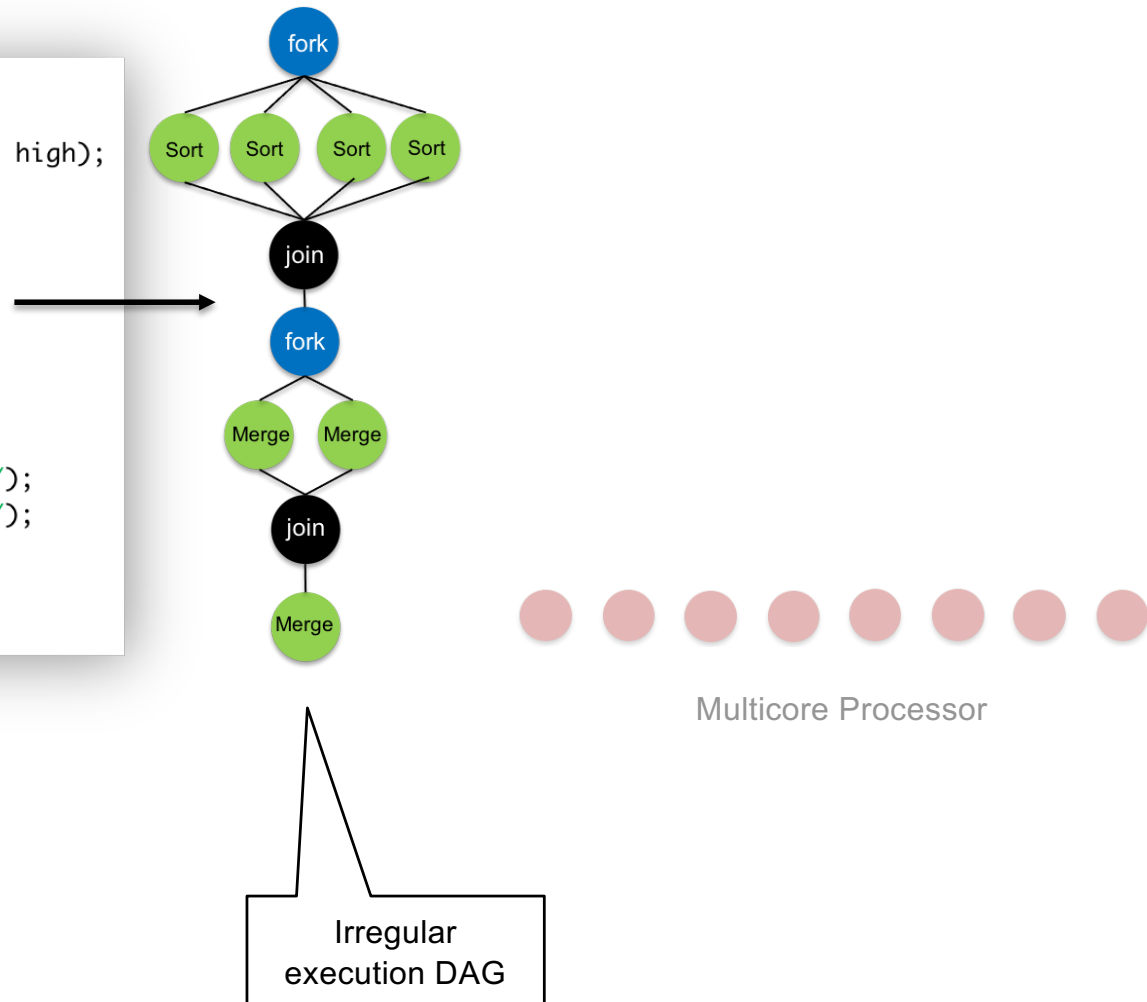
Serial elision
High productivity

Task Parallelism on Multicore Processors

```

1. int *A;
2. void Sort(int low, int high) {
3.   if((high-low)<LIMIT) return SeqSort(low, high);
4.   int Chunks=(high-low)/4;
5.   finish {
6.     async Sort(/*Chunk C1*/);
7.     async Sort(/*Chunk C2*/);
8.     async Sort(/*Chunk C3*/);
9.     async Sort(/*Chunk C4*/);
10.  }
11.  finish {
12.    async Merge(/*Chunk C1*/, /*Chunk C2*/);
13.    async Merge(/*Chunk C3*/, /*Chunk C4*/);
14.  }
15.  Merge(/*Chunk C12*/, /*Chunk C34*/);
16.}

```

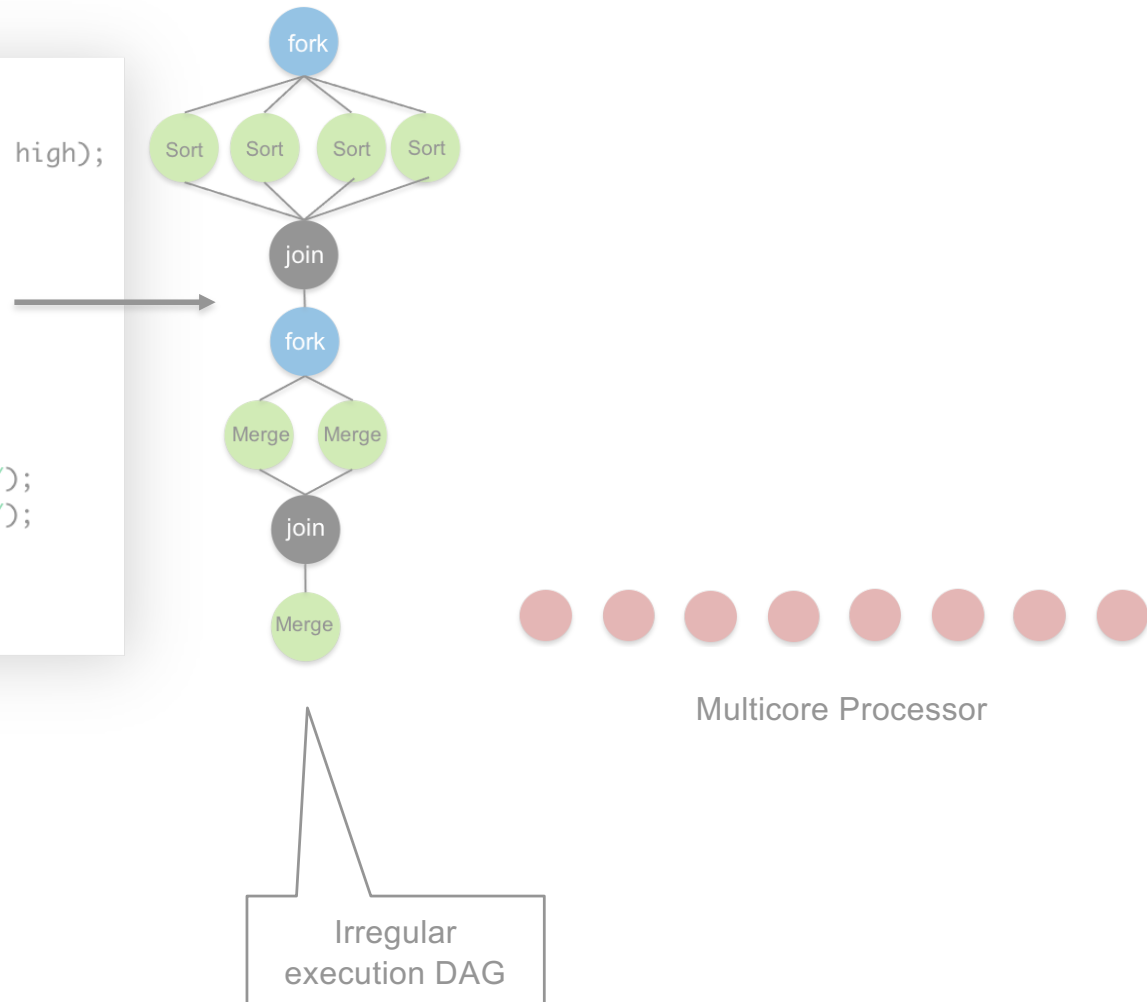


Task Parallelism on Multicore Processors

```

1. int *A;
2. void Sort(int low, int high) {
3.   if((high-low)<LIMIT) return SeqSort(low, high);
4.   int Chunks=(high-low)/4;
5.   finish {
6.     async Sort(/*Chunk C1*/);
7.     async Sort(/*Chunk C2*/);
8.     async Sort(/*Chunk C3*/);
9.     async Sort(/*Chunk C4*/);
10.  }
11.  finish {
12.    async Merge(/*Chunk C1*/, /*Chunk C2*/);
13.    async Merge(/*Chunk C3*/, /*Chunk C4*/);
14.  }
15.  Merge(/*Chunk C12*/, /*Chunk C34*/);
16.}
    
```

Serial elision
High productivity



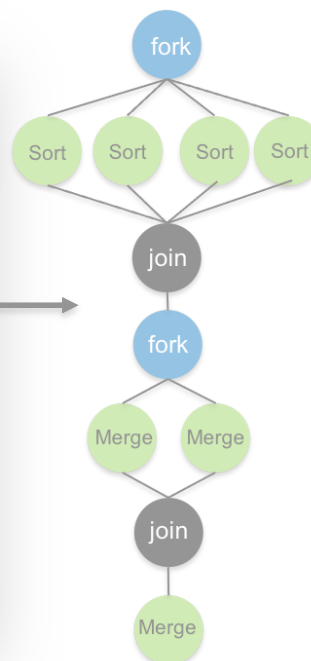
Task Parallelism on Multicore Processors

```

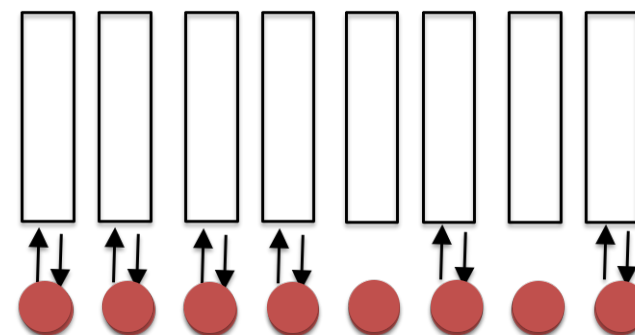
1. int *A;
2. void Sort(int low, int high) {
3.   if((high-low)<LIMIT) return SeqSort(low, high);
4.   int Chunks=(high-low)/4;
5.   finish {
6.     async Sort(/*Chunk C1*/);
7.     async Sort(/*Chunk C2*/);
8.     async Sort(/*Chunk C3*/);
9.     async Sort(/*Chunk C4*/);
10.  }
11.  finish {
12.    async Merge(/*Chunk C1*/, /*Chunk C2*/);
13.    async Merge(/*Chunk C3*/, /*Chunk C4*/);
14.  }
15.  Merge(/*Chunk C12*/, /*Chunk C34*/);
16.}

```

Serial elision
High productivity



Irregular
execution DAG



Multicore Processor

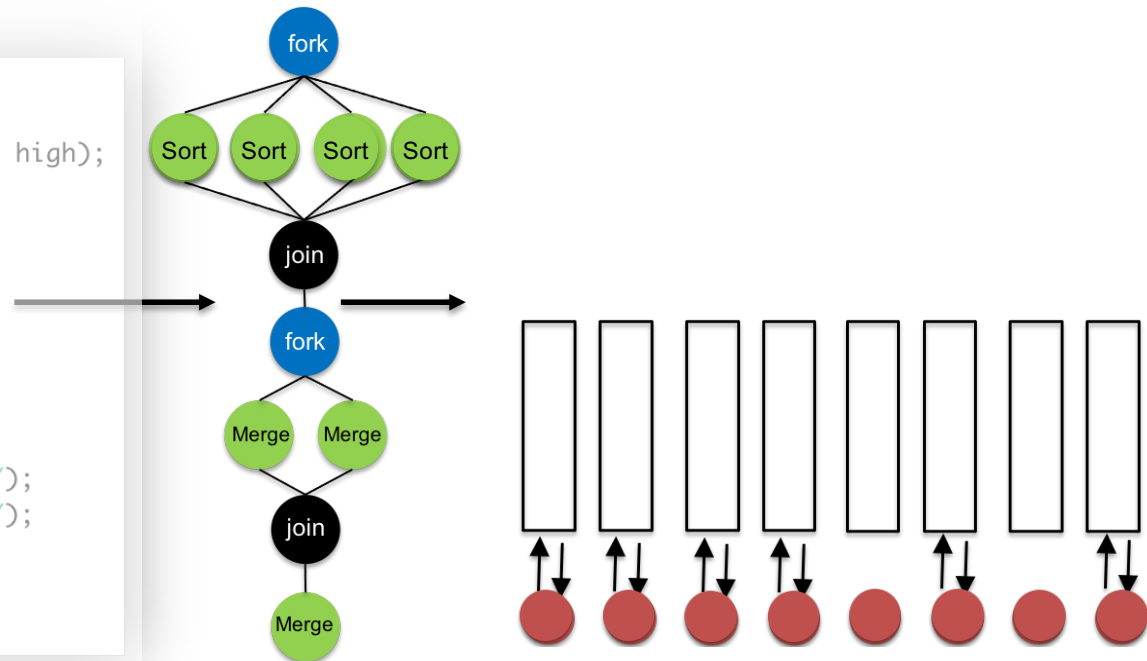
Task Parallelism on Multicore Processors

```

1. int *A;
2. void Sort(int low, int high) {
3.   if((high-low)<LIMIT) return SeqSort(low, high);
4.   int Chunks=(high-low)/4;
5.   finish {
6.     async Sort(/*Chunk C1*/);
7.     async Sort(/*Chunk C2*/);
8.     async Sort(/*Chunk C3*/);
9.     async Sort(/*Chunk C4*/);
10.  }
11.  finish {
12.    async Merge(/*Chunk C1*/, /*Chunk C2*/);
13.    async Merge(/*Chunk C3*/, /*Chunk C4*/);
14.  }
15.  Merge(/*Chunk C12*/, /*Chunk C34*/);
16.}

```

Serial elision
High productivity



Multicore Processor

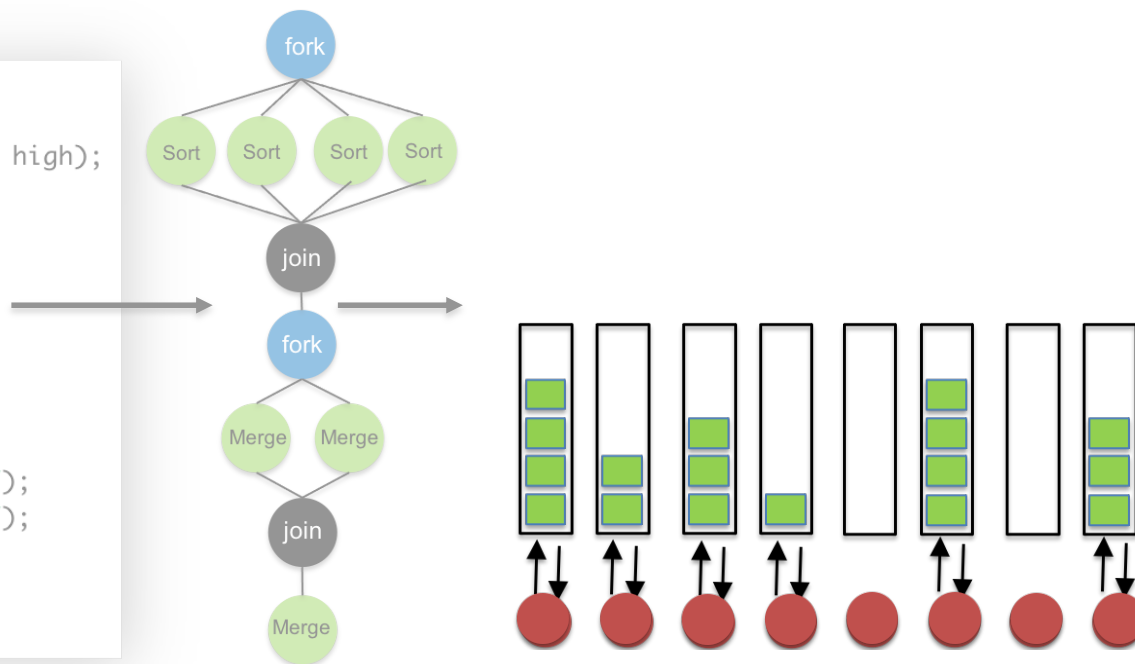
Irregular
execution DAG

Task Parallelism on Multicore Processors

```

1. int *A;
2. void Sort(int low, int high) {
3.   if((high-low)<LIMIT) return SeqSort(low, high);
4.   int Chunks=(high-low)/4;
5.   finish {
6.     async Sort(/*Chunk C1*/);
7.     async Sort(/*Chunk C2*/);
8.     async Sort(/*Chunk C3*/);
9.     async Sort(/*Chunk C4*/);
10.  }
11.  finish {
12.    async Merge(/*Chunk C1*/, /*Chunk C2*/);
13.    async Merge(/*Chunk C3*/, /*Chunk C4*/);
14.  }
15.  Merge(/*Chunk C12*/, /*Chunk C34*/);
16.}

```



Serial elision
High productivity

Irregular
execution DAG

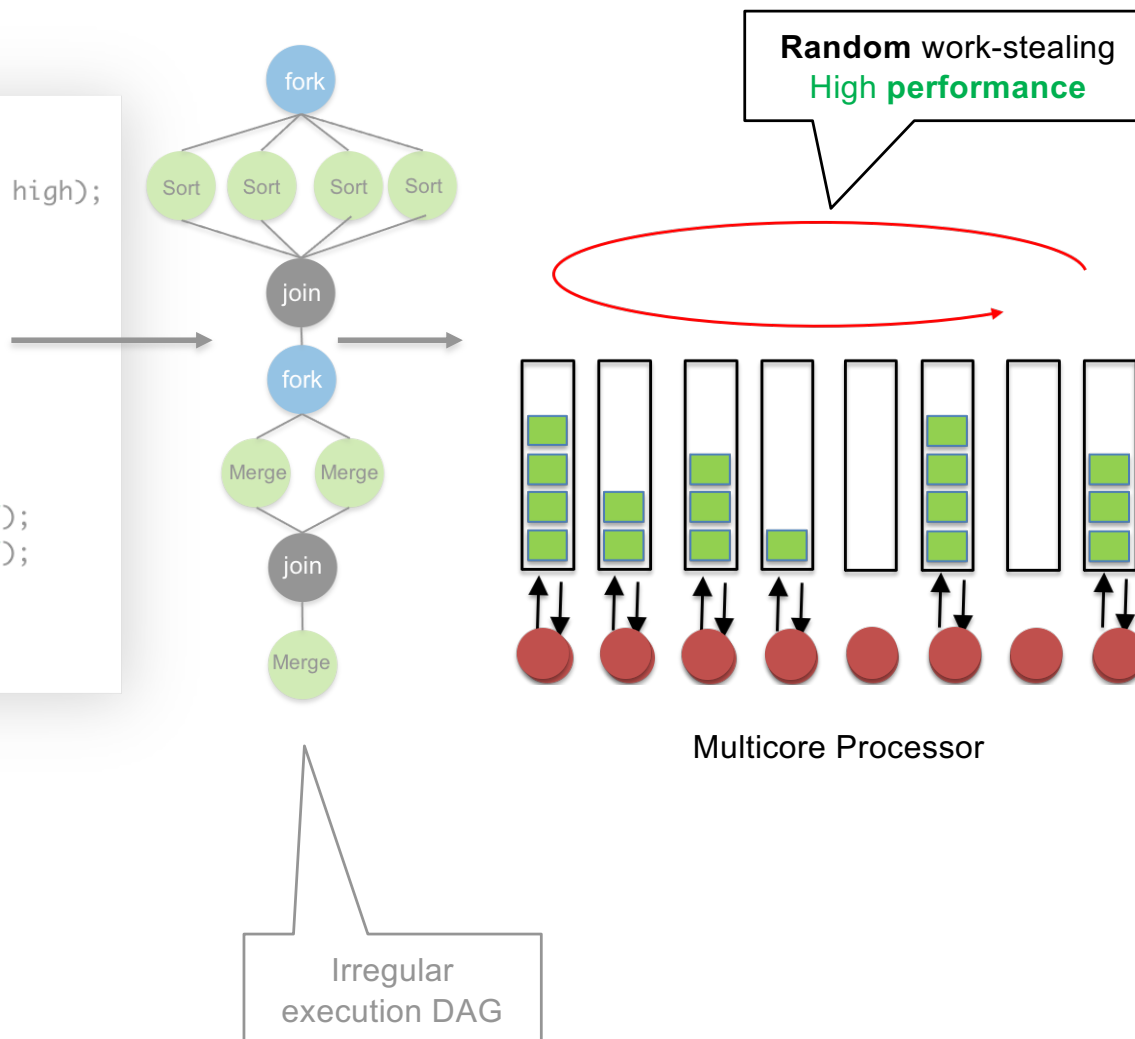
Task Parallelism on Multicore Processors

```

1. int *A;
2. void Sort(int low, int high) {
3.   if((high-low)<LIMIT) return SeqSort(low, high);
4.   int Chunks=(high-low)/4;
5.   finish {
6.     async Sort(/*Chunk C1*/);
7.     async Sort(/*Chunk C2*/);
8.     async Sort(/*Chunk C3*/);
9.     async Sort(/*Chunk C4*/);
10.  }
11.  finish {
12.    async Merge(/*Chunk C1*/, /*Chunk C2*/);
13.    async Merge(/*Chunk C3*/, /*Chunk C4*/);
14.  }
15.  Merge(/*Chunk C12*/, /*Chunk C34*/);
16.}

```

Serial elision
High productivity



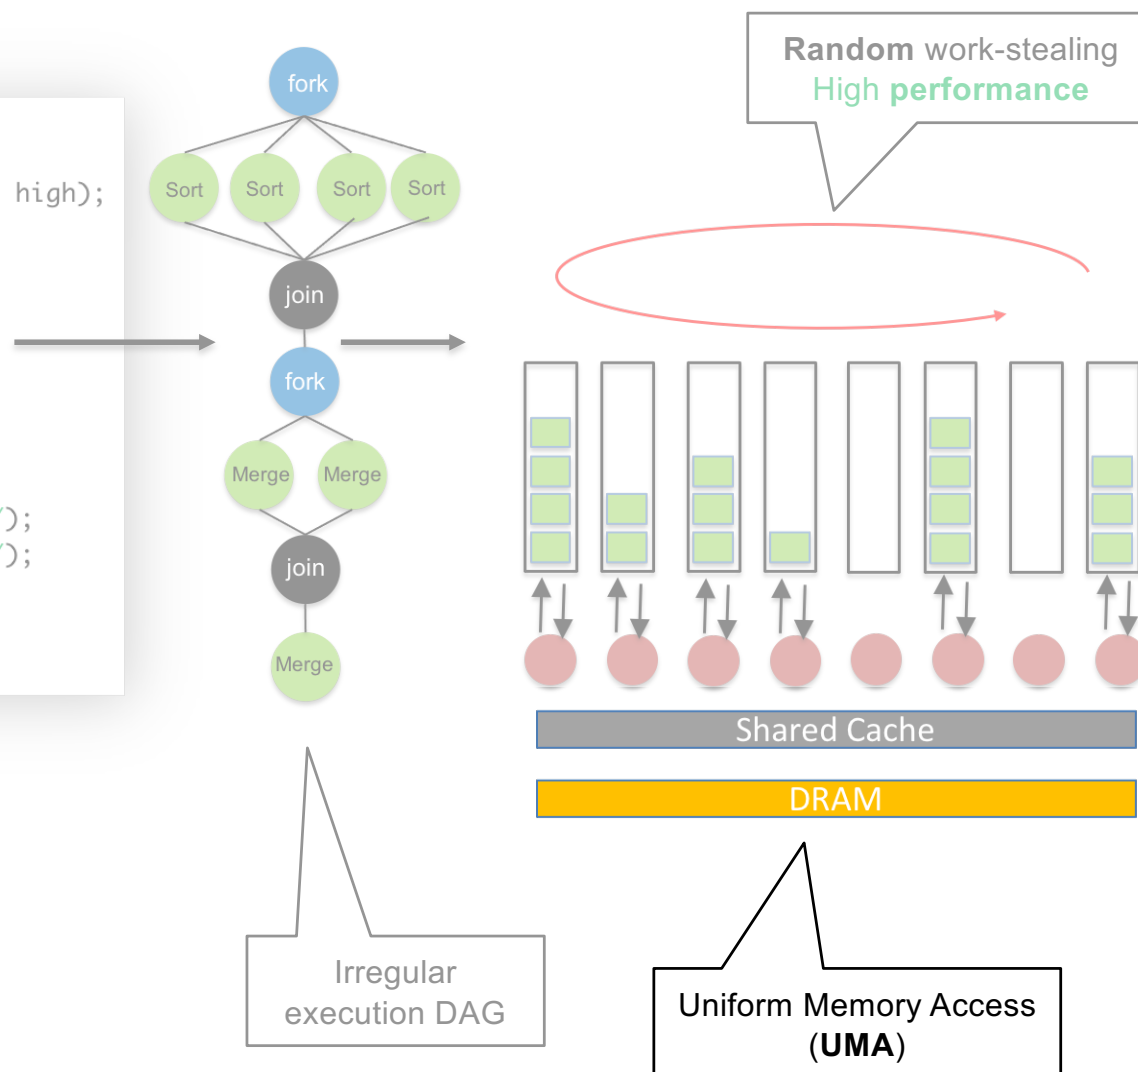
Task Parallelism on **UMA** Multicore Processors

```

1. int *A;
2. void Sort(int low, int high) {
3.   if((high-low)<LIMIT) return SeqSort(low, high);
4.   int Chunks=(high-low)/4;
5.   finish {
6.     async Sort(/*Chunk C1*/);
7.     async Sort(/*Chunk C2*/);
8.     async Sort(/*Chunk C3*/);
9.     async Sort(/*Chunk C4*/);
10.  }
11.  finish {
12.    async Merge(/*Chunk C1*/, /*Chunk C2*/);
13.    async Merge(/*Chunk C3*/, /*Chunk C4*/);
14.  }
15.  Merge(/*Chunk C12*/, /*Chunk C34*/);
16.}

```

Serial elision
High **productivity**

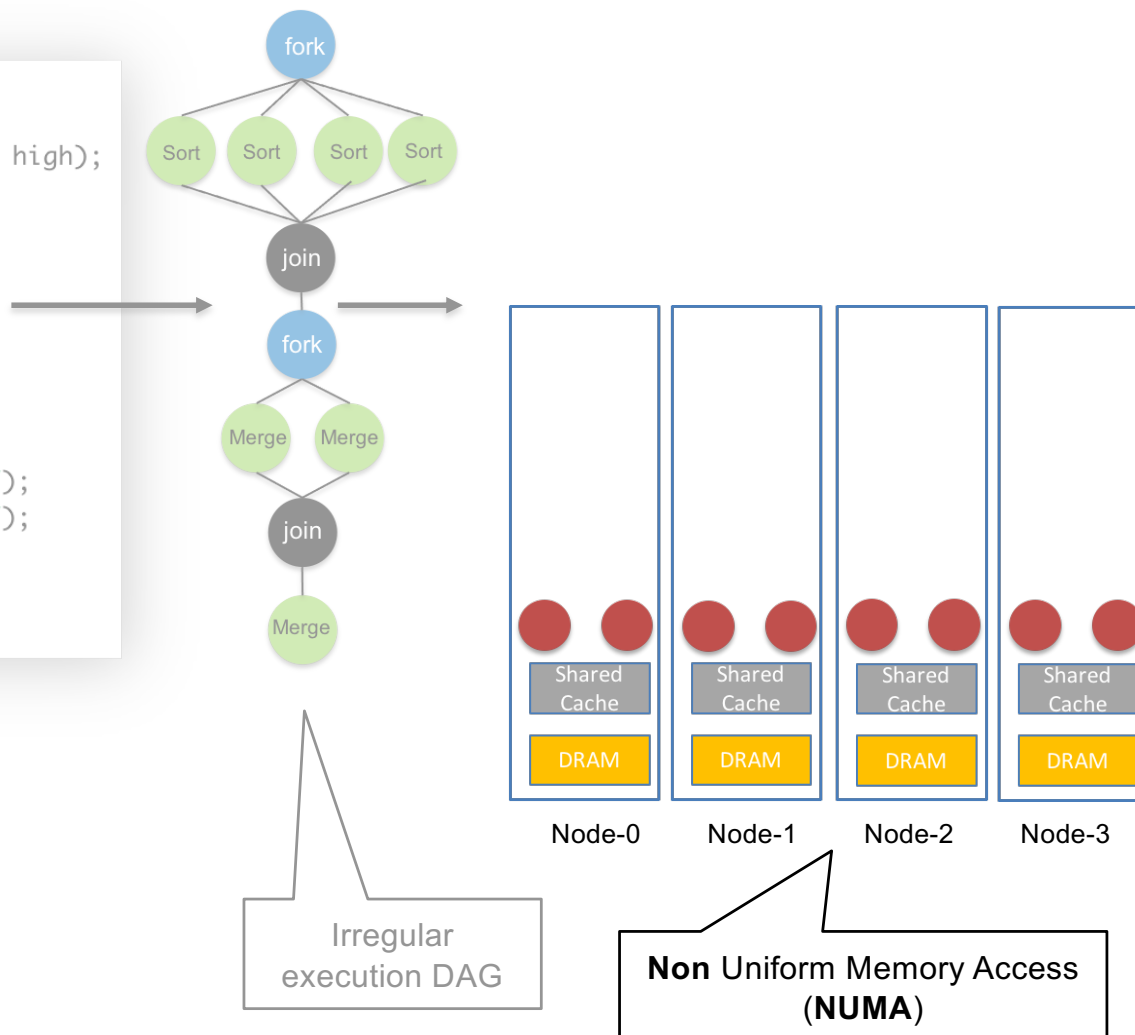


Task Parallelism on **NUMA** Multicore Processors

```

1. int *A;
2. void Sort(int low, int high) {
3.   if((high-low)<LIMIT) return SeqSort(low, high);
4.   int Chunks=(high-low)/4;
5.   finish {
6.     async Sort(/*Chunk C1*/);
7.     async Sort(/*Chunk C2*/);
8.     async Sort(/*Chunk C3*/);
9.     async Sort(/*Chunk C4*/);
10.  }
11.  finish {
12.    async Merge(/*Chunk C1*/, /*Chunk C2*/);
13.    async Merge(/*Chunk C3*/, /*Chunk C4*/);
14.  }
15.  Merge(/*Chunk C12*/, /*Chunk C34*/);
16.}

```



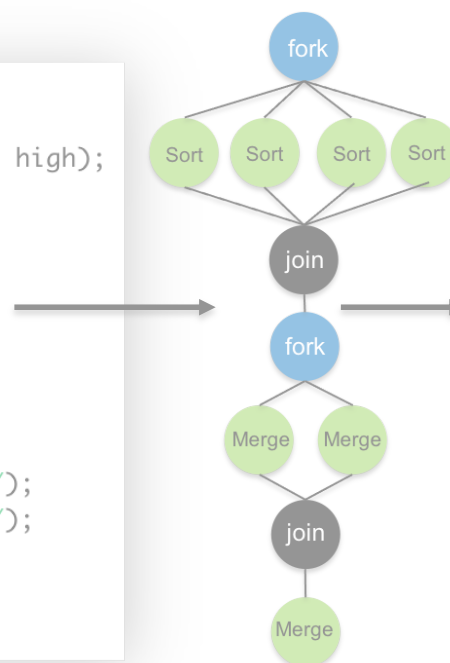
Task Parallelism on **NUMA** Multicore Processors

```

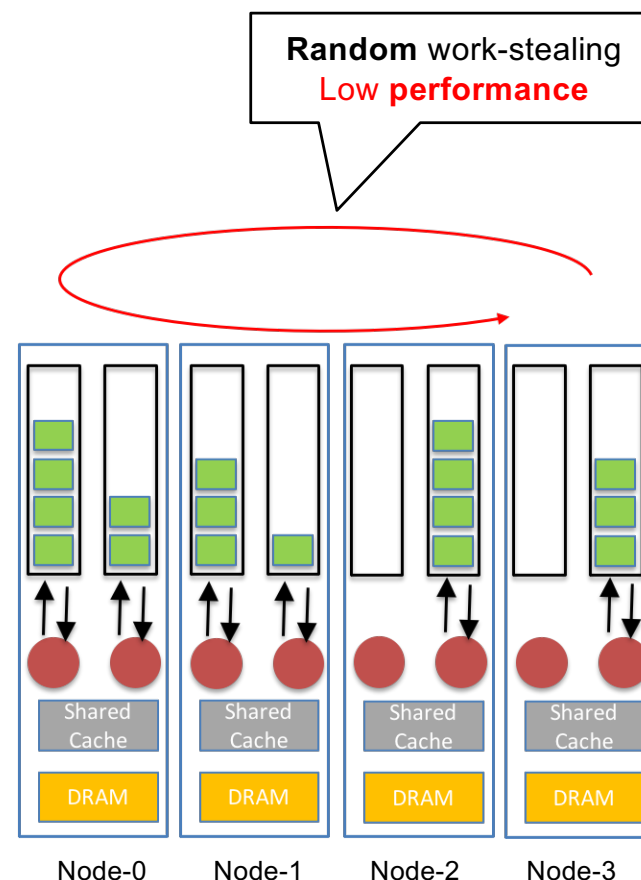
1. int *A;
2. void Sort(int low, int high) {
3.   if((high-low)<LIMIT) return SeqSort(low, high);
4.   int Chunks=(high-low)/4;
5.   finish {
6.     async Sort(/*Chunk C1*/);
7.     async Sort(/*Chunk C2*/);
8.     async Sort(/*Chunk C3*/);
9.     async Sort(/*Chunk C4*/);
10.  }
11.  finish {
12.    async Merge(/*Chunk C1*/, /*Chunk C2*/);
13.    async Merge(/*Chunk C3*/, /*Chunk C4*/);
14.  }
15.  Merge(/*Chunk C12*/, /*Chunk C34*/);
16.}

```

Serial elision
High productivity



Irregular
execution DAG



Non Uniform Memory Access
(NUMA)

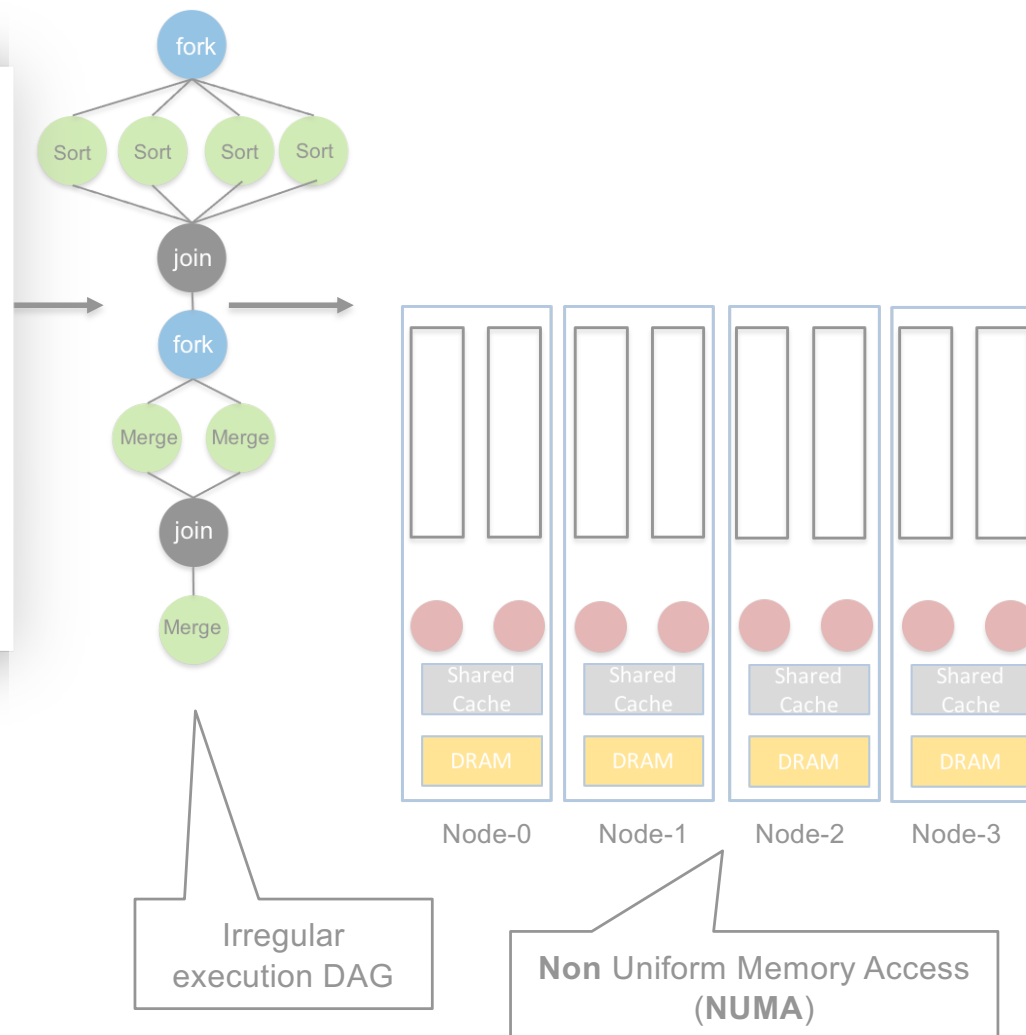
Task Parallelism on **NUMA** Multicore Processors

```

1. int *A;
2. void Sort(int low, int high) {
3.   if((high-low)<LIMIT) return SeqSort(low, high);
4.   int Chunks=(high-low)/4;
5.   finish {
6.     async_at(0) Sort(/*Chunk C1*/);
7.     async_at(1) Sort(/*Chunk C2*/);
8.     async_at(2) Sort(/*Chunk C3*/);
9.     async_at(3) Sort(/*Chunk C4*/);
10.  }
11.  finish {
12.    async_at(0) Merge(/*Chunk C1*/, /*Chunk C2*/);
13.    async_at(1) Merge(/*Chunk C3*/, /*Chunk C4*/);
14.  }
15.  Merge(/*Chunk C12*/, /*Chunk C34*/);
16.}

```

Serial elision
High productivity

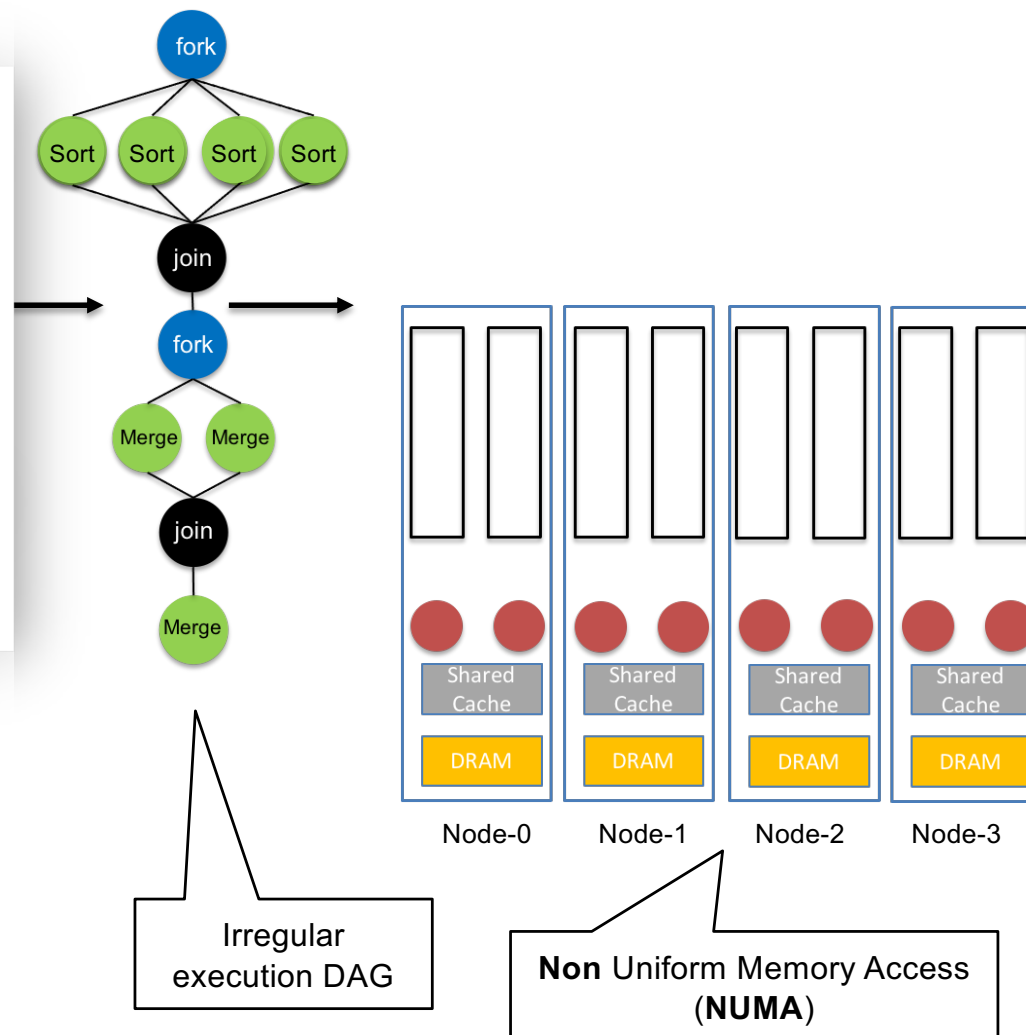


Task Parallelism on **NUMA** Multicore Processors

```

1. int *A;
2. void Sort(int low, int high) {
3.   if((high-low)<LIMIT) return SeqSort(low, high);
4.   int Chunks=(high-low)/4;
5.   finish {
6.     async_at(0) Sort(/*Chunk C1*/);
7.     async_at(1) Sort(/*Chunk C2*/);
8.     async_at(2) Sort(/*Chunk C3*/);
9.     async_at(3) Sort(/*Chunk C4*/);
10.  }
11.  finish {
12.    async_at(0) Merge(/*Chunk C1*/, /*Chunk C2*/);
13.    async_at(1) Merge(/*Chunk C3*/, /*Chunk C4*/);
14.  }
15.  Merge(/*Chunk C12*/, /*Chunk C34*/);
16.}
    
```

Serial elision
High **productivity**



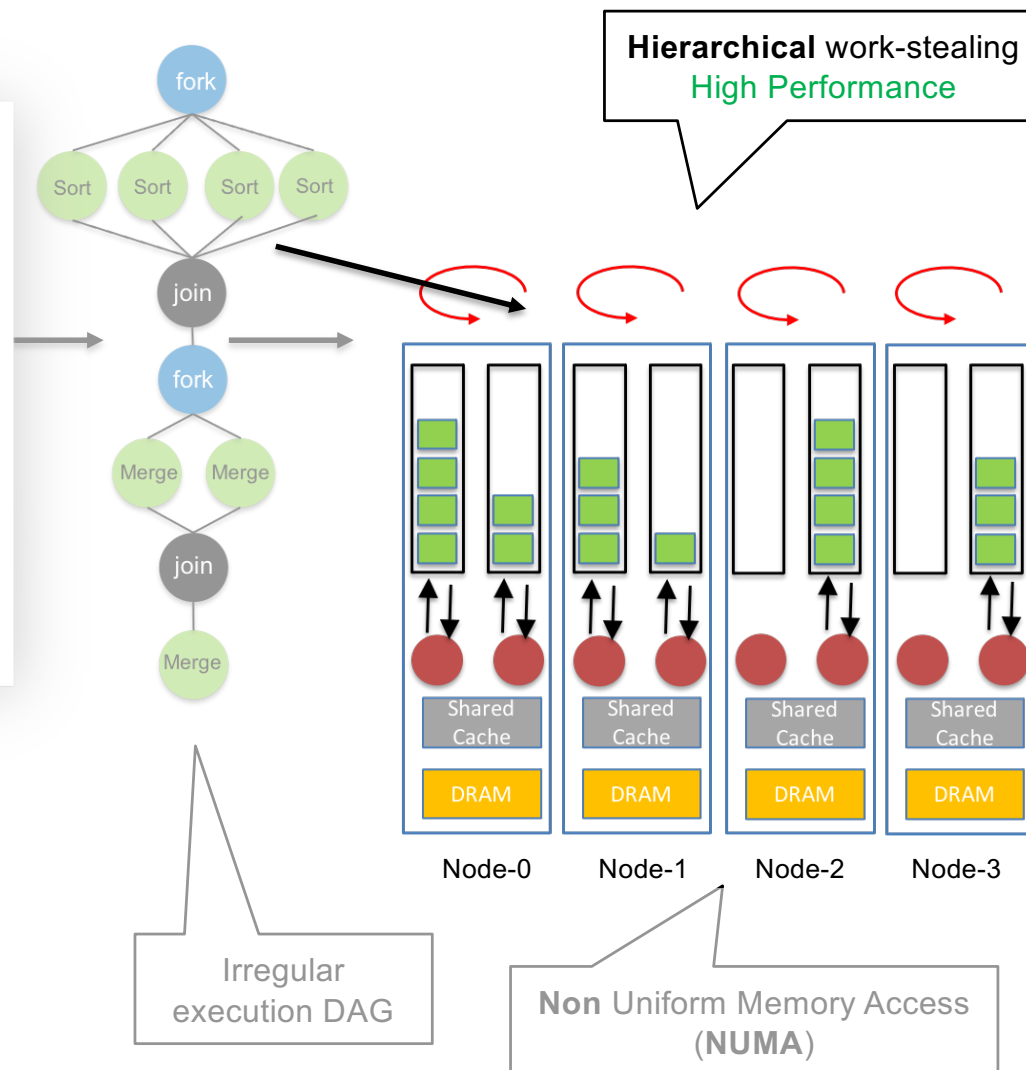
Task Parallelism on **NUMA** Multicore Processors

```

1. int *A;
2. void Sort(int low, int high) {
3.   if((high-low)<LIMIT) return SeqSort(low, high);
4.   int Chunks=(high-low)/4;
5.   finish {
6.     async_at(0) Sort(/*Chunk C1*/);
7.     async_at(1) Sort(/*Chunk C2*/);
8.     async_at(2) Sort(/*Chunk C3*/);
9.     async_at(3) Sort(/*Chunk C4*/);
10.  }
11.  finish {
12.    async_at(0) Merge(/*Chunk C1*/, /*Chunk C2*/);
13.    async_at(1) Merge(/*Chunk C3*/, /*Chunk C4*/);
14.  }
15.  Merge(/*Chunk C12*/, /*Chunk C34*/);
16.}

```

Serial elision
High productivity

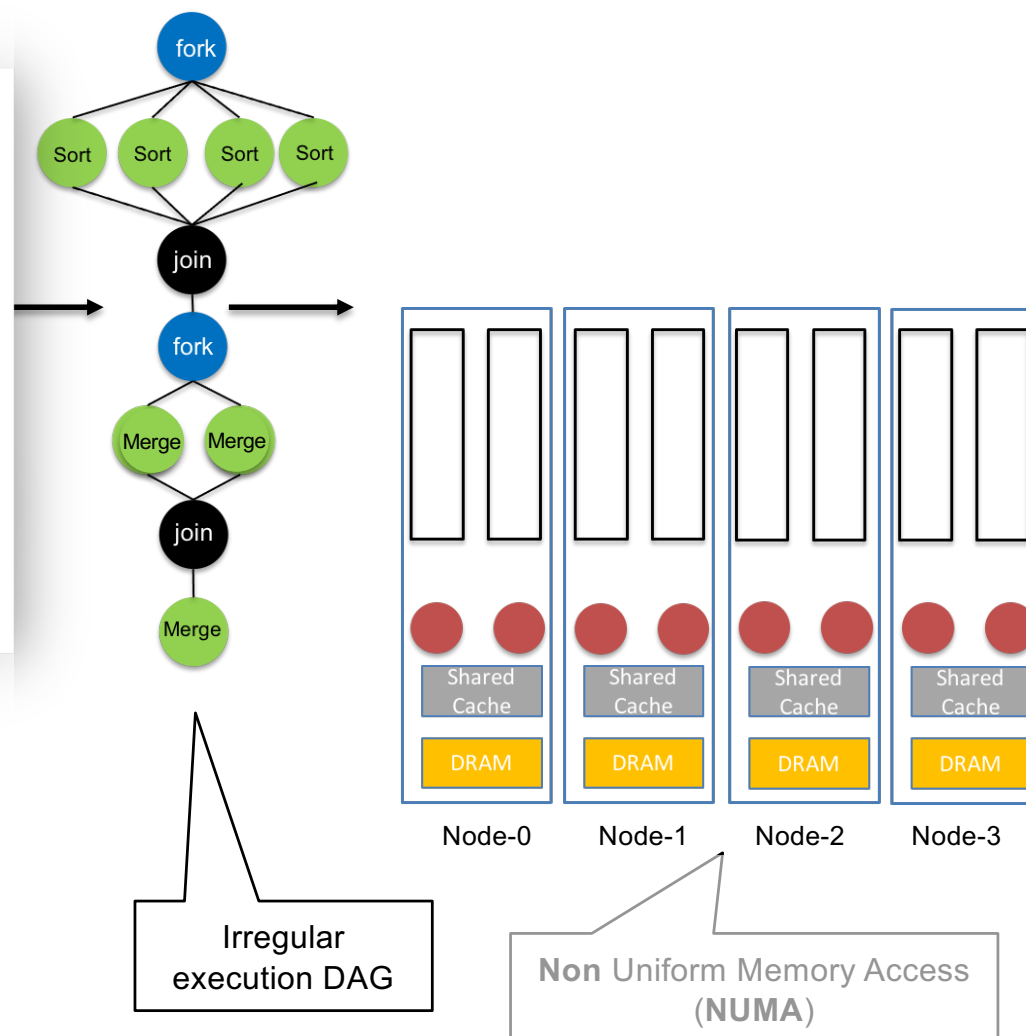


Task Parallelism on **NUMA** Multicore Processors

```

1. int *A;
2. void Sort(int low, int high) {
3.   if((high-low)<LIMIT) return SeqSort(low, high);
4.   int Chunks=(high-low)/4;
5.   finish {
6.     async_at(0) Sort(/*Chunk C1*/);
7.     async_at(1) Sort(/*Chunk C2*/);
8.     async_at(2) Sort(/*Chunk C3*/);
9.     async_at(3) Sort(/*Chunk C4*/);
10.  }
11.  finish {
12.    async_at(0) Merge(/*Chunk C1*/, /*Chunk C2*/);
13.    async_at(1) Merge(/*Chunk C3*/, /*Chunk C4*/);
14.  }
15.  Merge(/*Chunk C12*/, /*Chunk C34*/);
16.}
    
```

Serial elision
High productivity



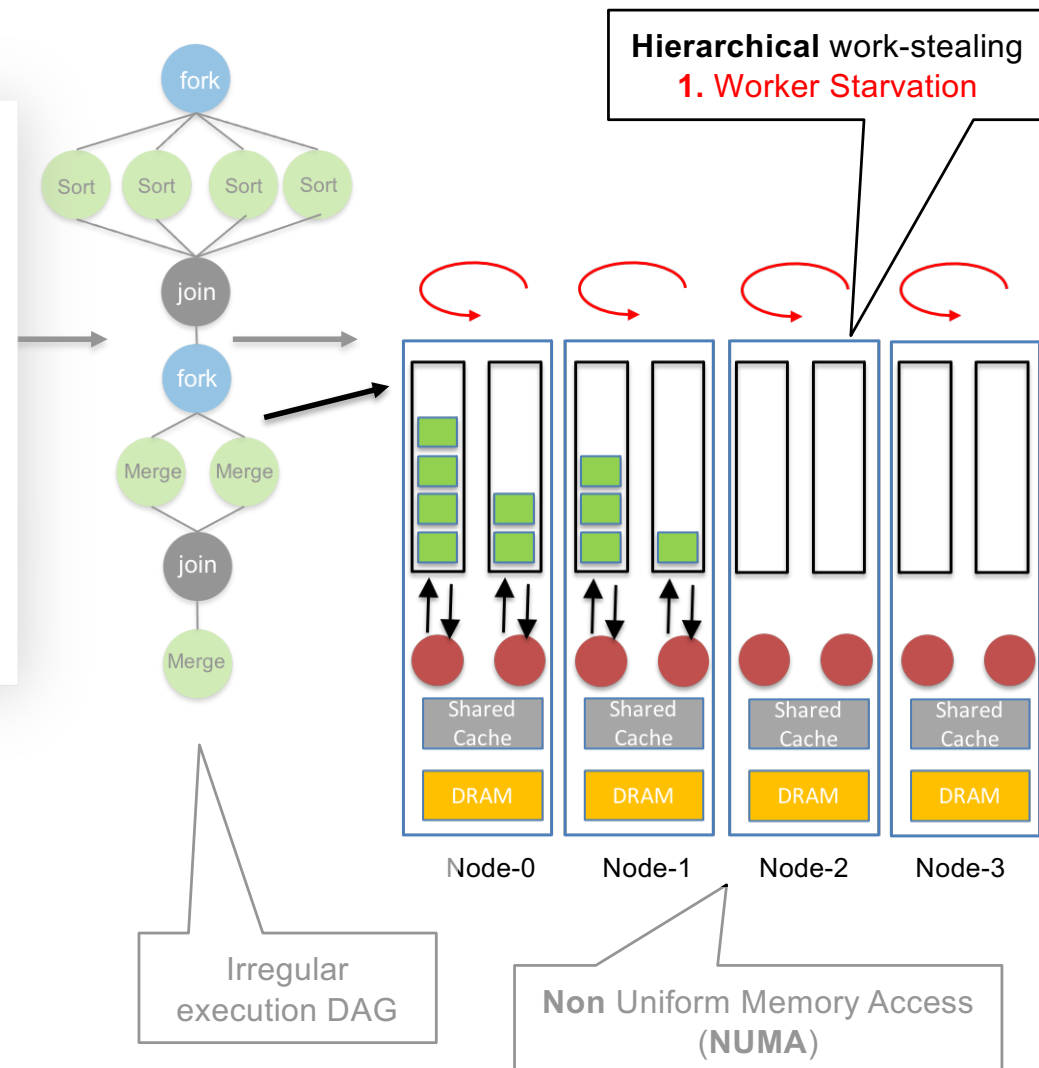
Task Parallelism on **NUMA** Multicore Processors

```

1. int *A;
2. void Sort(int low, int high) {
3.   if((high-low)<LIMIT) return SeqSort(low, high);
4.   int Chunks=(high-low)/4;
5.   finish {
6.     async_at(0) Sort(/*Chunk C1*/);
7.     async_at(1) Sort(/*Chunk C2*/);
8.     async_at(2) Sort(/*Chunk C3*/);
9.     async_at(3) Sort(/*Chunk C4*/);
10.  }
11.  finish {
12.    async_at(0) Merge(/*Chunk C1*/, /*Chunk C2*/);
13.    async_at(1) Merge(/*Chunk C3*/, /*Chunk C4*/);
14.  }
15.  Merge(/*Chunk C12*/, /*Chunk C34*/);
16.}

```

Serial elision
High productivity



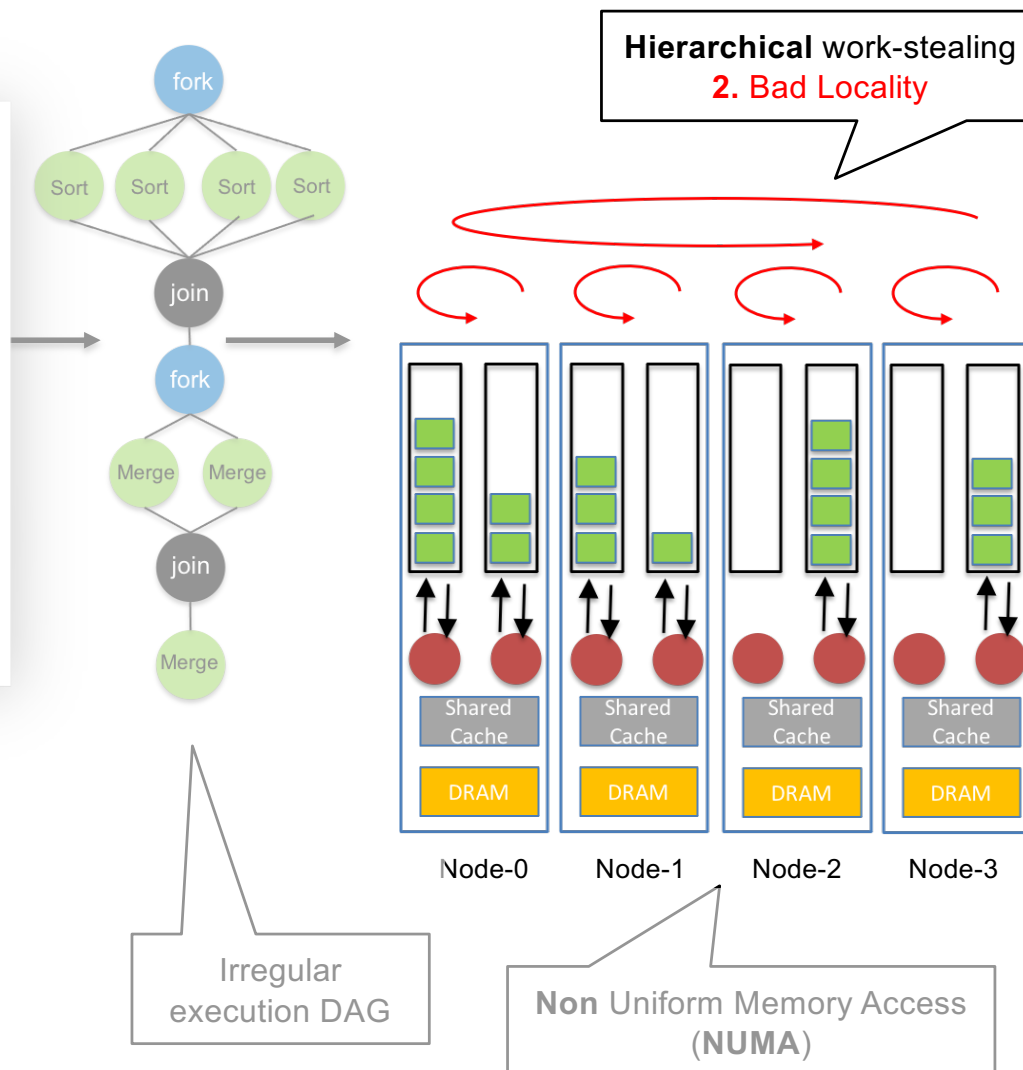
Task Parallelism on **NUMA** Multicore Processors

```

1. int *A;
2. void Sort(int low, int high) {
3.   if((high-low)<LIMIT) return SeqSort(low, high);
4.   int Chunks=(high-low)/4;
5.   finish {
6.     async_at(0) Sort(/*Chunk C1*/);
7.     async_at(1) Sort(/*Chunk C2*/);
8.     async_at(2) Sort(/*Chunk C3*/);
9.     async_at(3) Sort(/*Chunk C4*/);
10.  }
11.  finish {
12.    async_at(0) Merge(/*Chunk C1*/, /*Chunk C2*/);
13.    async_at(1) Merge(/*Chunk C3*/, /*Chunk C4*/);
14.  }
15.  Merge(/*Chunk C12*/, /*Chunk C34*/);
16.}

```

Serial elision
High productivity



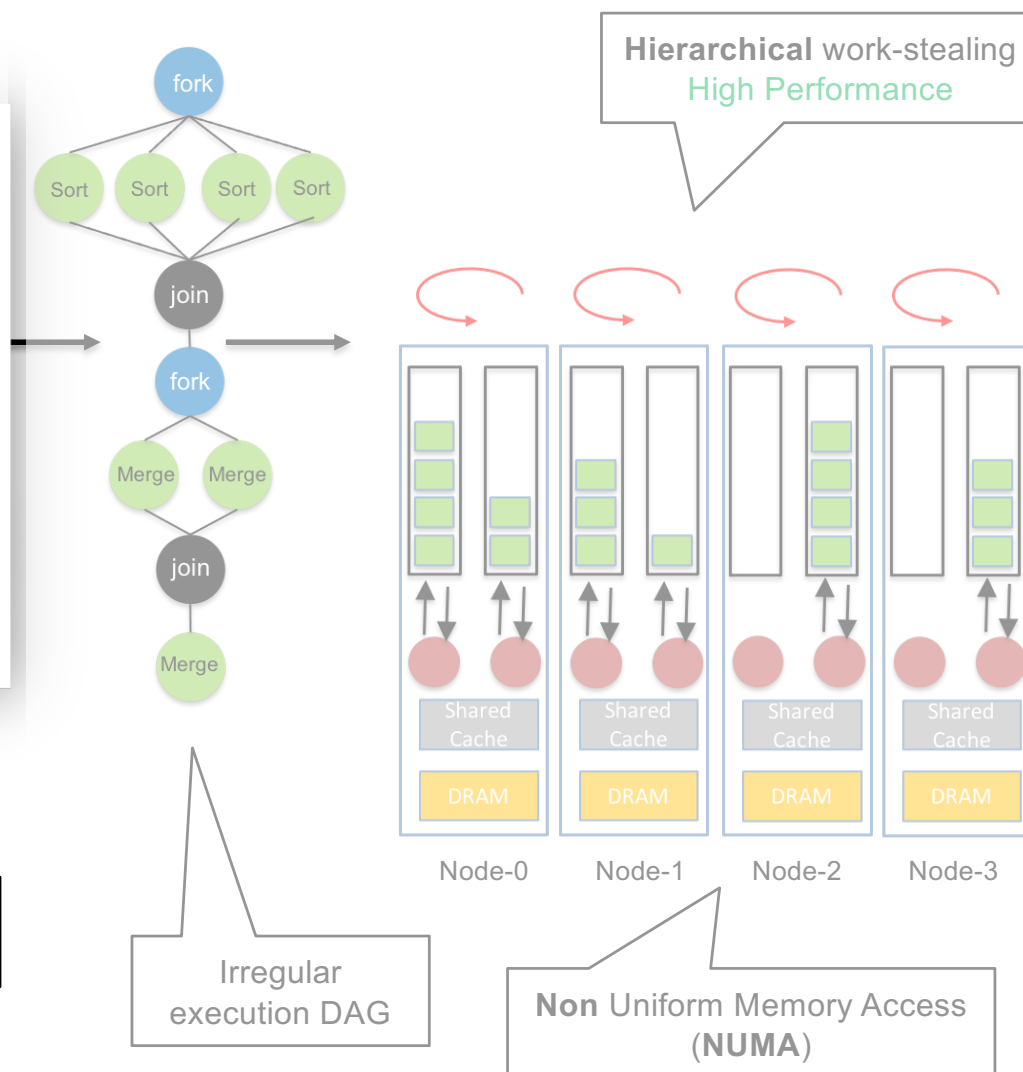
Task Parallelism on **NUMA** Multicore Processors

```

1. int *A;
2. void Sort(int low, int high) {
3.   if((high-low)<LIMIT) return SeqSort(low, high);
4.   int Chunks=(high-low)/4;
5.   finish {
6.     async_at(0) Sort(/*Chunk C1*/);
7.     async_at(1) Sort(/*Chunk C2*/);
8.     async_at(2) Sort(/*Chunk C3*/);
9.     async_at(3) Sort(/*Chunk C4*/);
10.  }
11.  finish {
12.    async_at(0) Merge(/*Chunk C1*/, /*Chunk C2*/);
13.    async_at(1) Merge(/*Chunk C3*/, /*Chunk C4*/);
14.  }
15.  Merge(/*Chunk C12*/, /*Chunk C34*/);
16.}
    
```

Modify "merge" into
four subtask

3. No Serial elision



Library for Task Parallelism & Work-Stealing over NUMA Processors

```

1. int *A;
2. void Sort(int low, int high) {
3.     if((high-low)<LIMIT) return SeqSort(low, high);
4.     int Chunks=(high-low)/4;
5.     finish {
6.         async_hinted (A, C1_start, C1_end) Sort(/*Chunk C1*/);
7.         async_hinted (A, C2_start, C2_end) Sort(/*Chunk C2*/);
8.         async_hinted (A, C3_start, C3_end) Sort(/*Chunk C3*/);
9.         async_hinted (A, C4_start, C4_end) Sort(/*Chunk C4*/);
10.    }
11.    finish {
12.        async_hinted (A, C1_start, C2_end) Merge(/*Chunk C1*/, /*Chunk C2*/);
13.        async_hinted (A, C3_start, C4_end) Merge(/*Chunk C3*/, /*Chunk C4*/);
14.    }
15.    Merge(/*Chunk C12*/, /*Chunk C34*/);
16.}

```

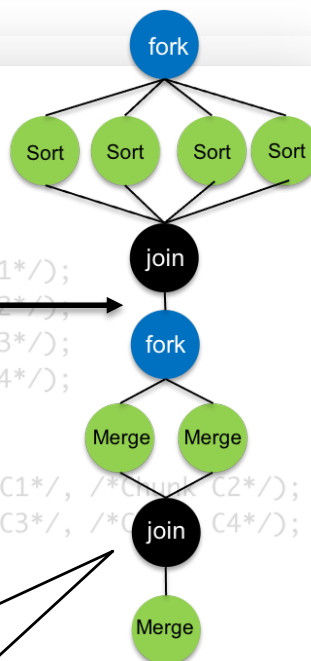
1. Supports Serial elision

(except for using NUMA aware malloc/free)

Library for Task Parallelism & Work-Stealing over NUMA Processors

```

1. int *A;
2. void Sort(int low, int high) {
3.   if((high-low)<LIMIT) return SeqSort(low, high);
4.   int Chunks=(high-low)/4;
5.   finish {
6.     async_hinted (A, C1_start, C1_end) Sort(/*Chunk C1*/);
7.     async_hinted (A, C2_start, C2_end) Sort(/*Chunk C2*/);
8.     async_hinted (A, C3_start, C3_end) Sort(/*Chunk C3*/);
9.     async_hinted (A, C4_start, C4_end) Sort(/*Chunk C4*/);
10.  }
11.  finish {
12.    async_hinted (A, C1_start, C2_end) Merge(/*Chunk C1*/, /*Chunk C2*/);
13.    async_hinted (A, C3_start, C4_end) Merge(/*Chunk C3*/, /*Chunk C4*/);
14.  }
15.  Merge(/*Chunk C12*/, /*Chunk C34*/);
16.}
    
```



2. Supports Irregular execution DAG

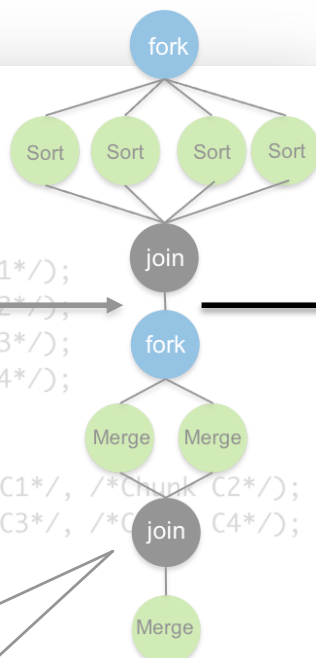
1. Supports Serial elision

(except for using NUMA aware malloc/free)

Library for Task Parallelism & Work-Stealing over NUMA Processors

```

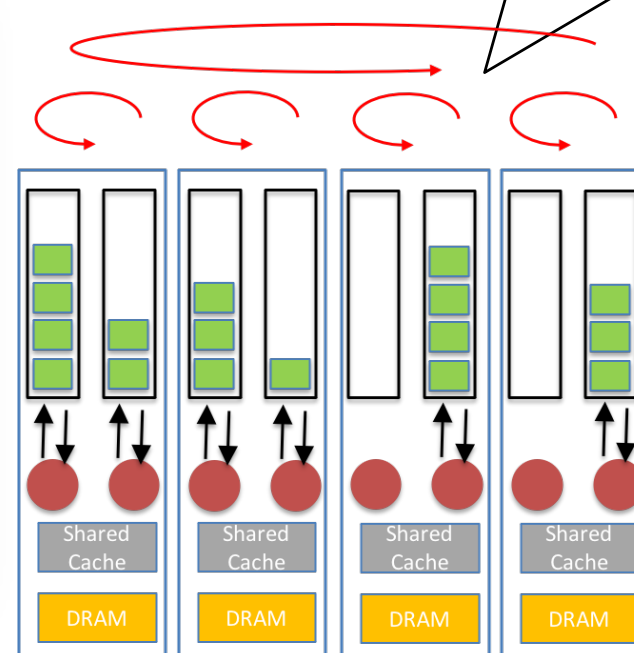
1. int *A;
2. void Sort(int low, int high) {
3.   if((high-low)<LIMIT) return SeqSort(low, high);
4.   int Chunks=(high-low)/4;
5.   finish {
6.     async_hinted (A, C1_start, C1_end) Sort(/*Chunk C1*/);
7.     async_hinted (A, C2_start, C2_end) Sort(/*Chunk C2*/);
8.     async_hinted (A, C3_start, C3_end) Sort(/*Chunk C3*/);
9.     async_hinted (A, C4_start, C4_end) Sort(/*Chunk C4*/);
10.  }
11. finish {
12.   async_hinted (A, C1_start, C2_end) Merge(/*Chunk C1*/, /*Chunk C2*/);
13.   async_hinted (A, C3_start, C4_end) Merge(/*Chunk C3*/, /*Chunk C4*/);
14. }
15. Merge(/*Chunk C12*/, /*Chunk C34*/);
16.}
    
```



2. Supports Irregular execution DAG

1. Supports Serial elision
(except for using NUMA aware malloc/free)

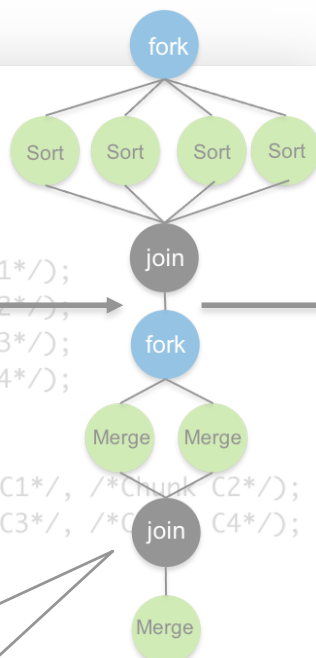
Hierarchical work-stealing
3. Improves Locality
4. Removes Starvation



Library for Task Parallelism & Work-Stealing over NUMA Processors

```

1. int *A;
2. void Sort(int low, int high) {
3.   if((high-low)<LIMIT) return SeqSort(low, high);
4.   int Chunks=(high-low)/4;
5.   finish {
6.     async_hinted (A, C1_start, C1_end) Sort(/*Chunk C1*/);
7.     async_hinted (A, C2_start, C2_end) Sort(/*Chunk C2*/);
8.     async_hinted (A, C3_start, C3_end) Sort(/*Chunk C3*/);
9.     async_hinted (A, C4_start, C4_end) Sort(/*Chunk C4*/);
10.  }
11.  finish {
12.    async_hinted (A, C1_start, C2_end) Merge(/*Chunk C1*/, /*Chunk C2*/);
13.    async_hinted (A, C3_start, C4_end) Merge(/*Chunk C3*/, /*Chunk C4*/);
14.  }
15.  Merge(/*Chunk C12*/, /*Chunk C34*/);
16.}
    
```

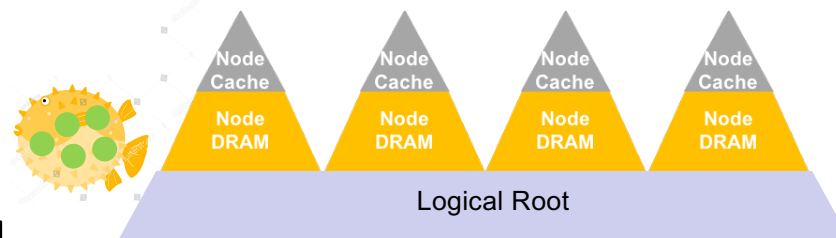
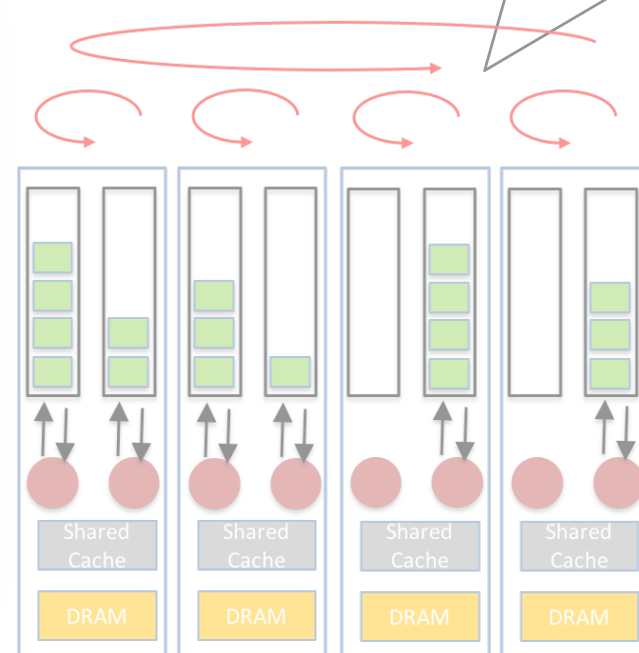


2. Supports Irregular execution DAG

1. Supports Serial elision
(except for using NUMA aware malloc/free)

5. Hierarchical Elastic Tasks
(further improves the locality)

Hierarchical work-stealing
3. Improves Locality
4. Removes Starvation

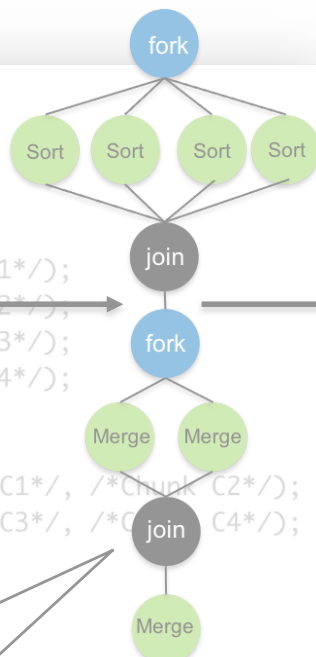


NUMA Memory Hierarchy

Library for Task Parallelism & Work-Stealing over NUMA Processors

```

1. int *A;
2. void Sort(int low, int high) {
3.   if((high-low)<LIMIT) return SeqSort(low, high);
4.   int Chunks=(high-low)/4;
5.   finish {
6.     async_hinted (A, C1_start, C1_end) Sort(/*Chunk C1*/);
7.     async_hinted (A, C2_start, C2_end) Sort(/*Chunk C2*/);
8.     async_hinted (A, C3_start, C3_end) Sort(/*Chunk C3*/);
9.     async_hinted (A, C4_start, C4_end) Sort(/*Chunk C4*/);
10.  }
11. finish {
12.   async_hinted (A, C1_start, C2_end) Merge(/*Chunk C1*/, /*Chunk C2*/);
13.   async_hinted (A, C3_start, C4_end) Merge(/*Chunk C3*/, /*Chunk C4*/);
14. }
15. Merge(/*Chunk C12*/, /*Chunk C34*/);
16.}
    
```

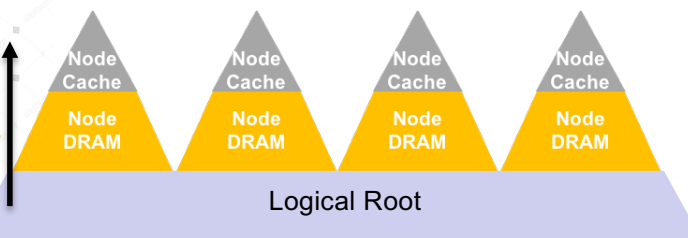
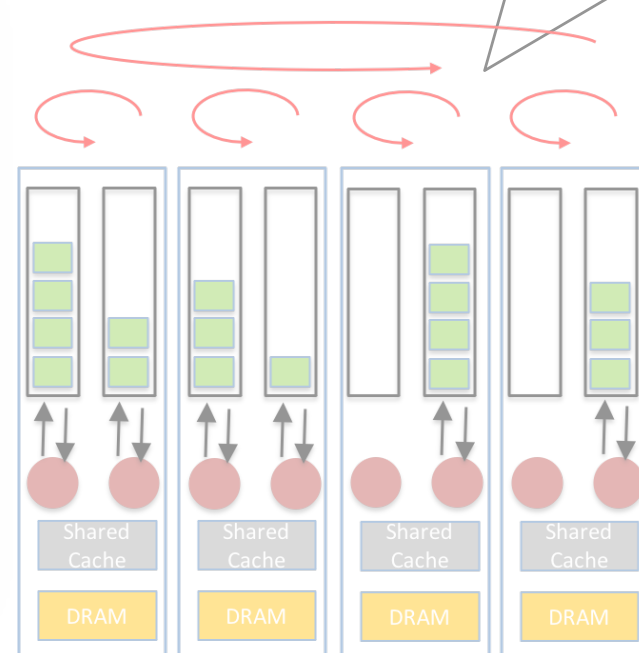


2. Supports Irregular execution DAG

1. Supports Serial elision
(except for using NUMA aware malloc/free)

5. Hierarchical Elastic Tasks
(further improves the locality)

Hierarchical work-stealing
3. Improves Locality
4. Removes Starvation

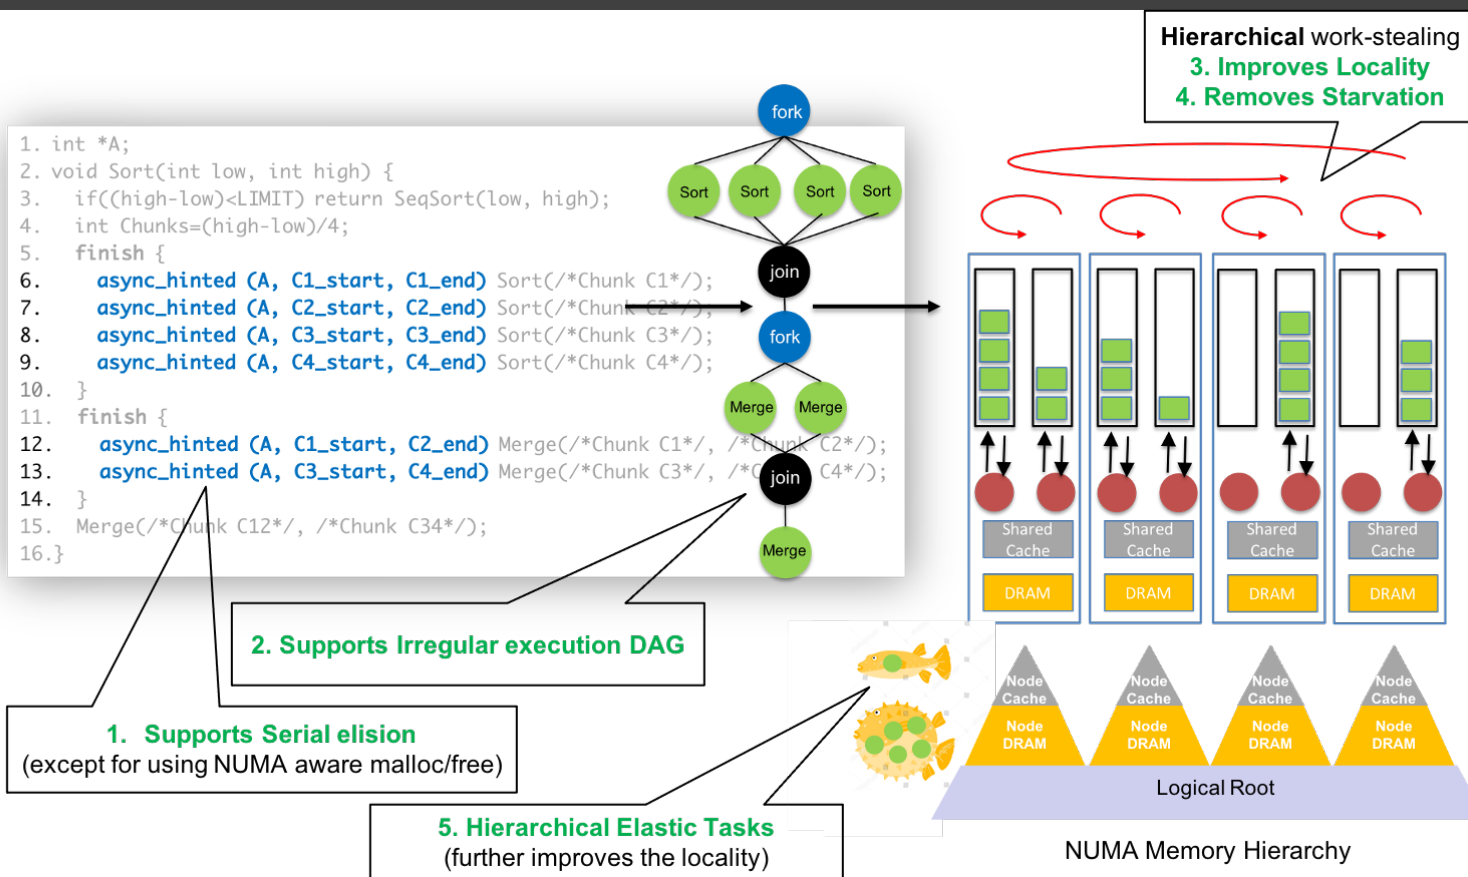


NUMA Memory Hierarchy

Performance Analysis on AMD EPYC 7551



Executing summary for **seven** recursive benchmarks with regular/irregular DAG on a **32-core** processor with **four** NUMA nodes



Artifact

