

Optimized Distributed Work-Stealing

Vivek Kumar¹, Karthik Murthy¹, Vivek Sarkar¹, Yili Zheng²


1 Rice University

2 Lawrence Berkeley National Laboratory




Multicore Nodes in Supercomputers

Cores/Socket System Share in Top500



November **2006**



June **2016**

Graph plotted using the data obtained from <https://www.top500.org/statistics/list/>



Problem Statement

Productivity and Performance Challenge

- Productivity
 - Several existing APIs for scientific computing
 - Hard to parallelize complex irregular computations using existing APIs
 - Ideal candidate for runtime based global load-balancing
- Performance on multicore nodes
 - Using a process per core (e.g., MPI everywhere) on a node not scalable
 - Hybrid programming using thread pool per node
 - How to design a high performance implementation of global load-balancing



Contributions

- ✓ Library-based API in a PGAS library to express irregular computations

C++11 lambda function based API that provides serial elision

- ✓ Novel implementation of distributed work-stealing

That introduces a new victim selection policy that avoid all inter-node failed steals

- ✓ Detailed performance study


That demonstrates the benefit using scaling irregular applications up to 12k cores of Edison supercomputer

- ✓ Results

That shows that our approach delivers performance benefits up to 7%



Load Balancing using Work-Stealing



Work-stealing in a thread pool


- Thread pool (**intra-node**) based implementations perform stealing using low overhead CAS operations



Motivating Analysis

Distributed Work-Stealing

- **Inter-node** steals are much costlier than **intra-node** steals



Intra-node
steals

Inter-node
steals

Motivating Analysis

Failed Steal Attempts


- Thief fails to steal a task from victim

Inter-node failed steals are more costly than **intra-node** steals


Chances to fail with same victim multiple times

Intra-node failed steals

Inter-node failed steals



Our Approach



One process with a thread pool at each node

- Use HabaneroUPC++ PGAS library for multicore cluster [Kumar et. al., PGAS 2014]
 - Several asynchronous tasking APIs
- Provide a programming model to express irregular computation
- Implement a high performance distributed work-stealing runtime that **completely removes** all inter-node failed steal attempts



HabaneroUPC++ Programming Model

```
asyncAny ( [=] {  
    irregular_computation();  
}); //distributed work-stealing
```

- C++11 lambda-function based API
- Provides **serial elision** and improves productivity



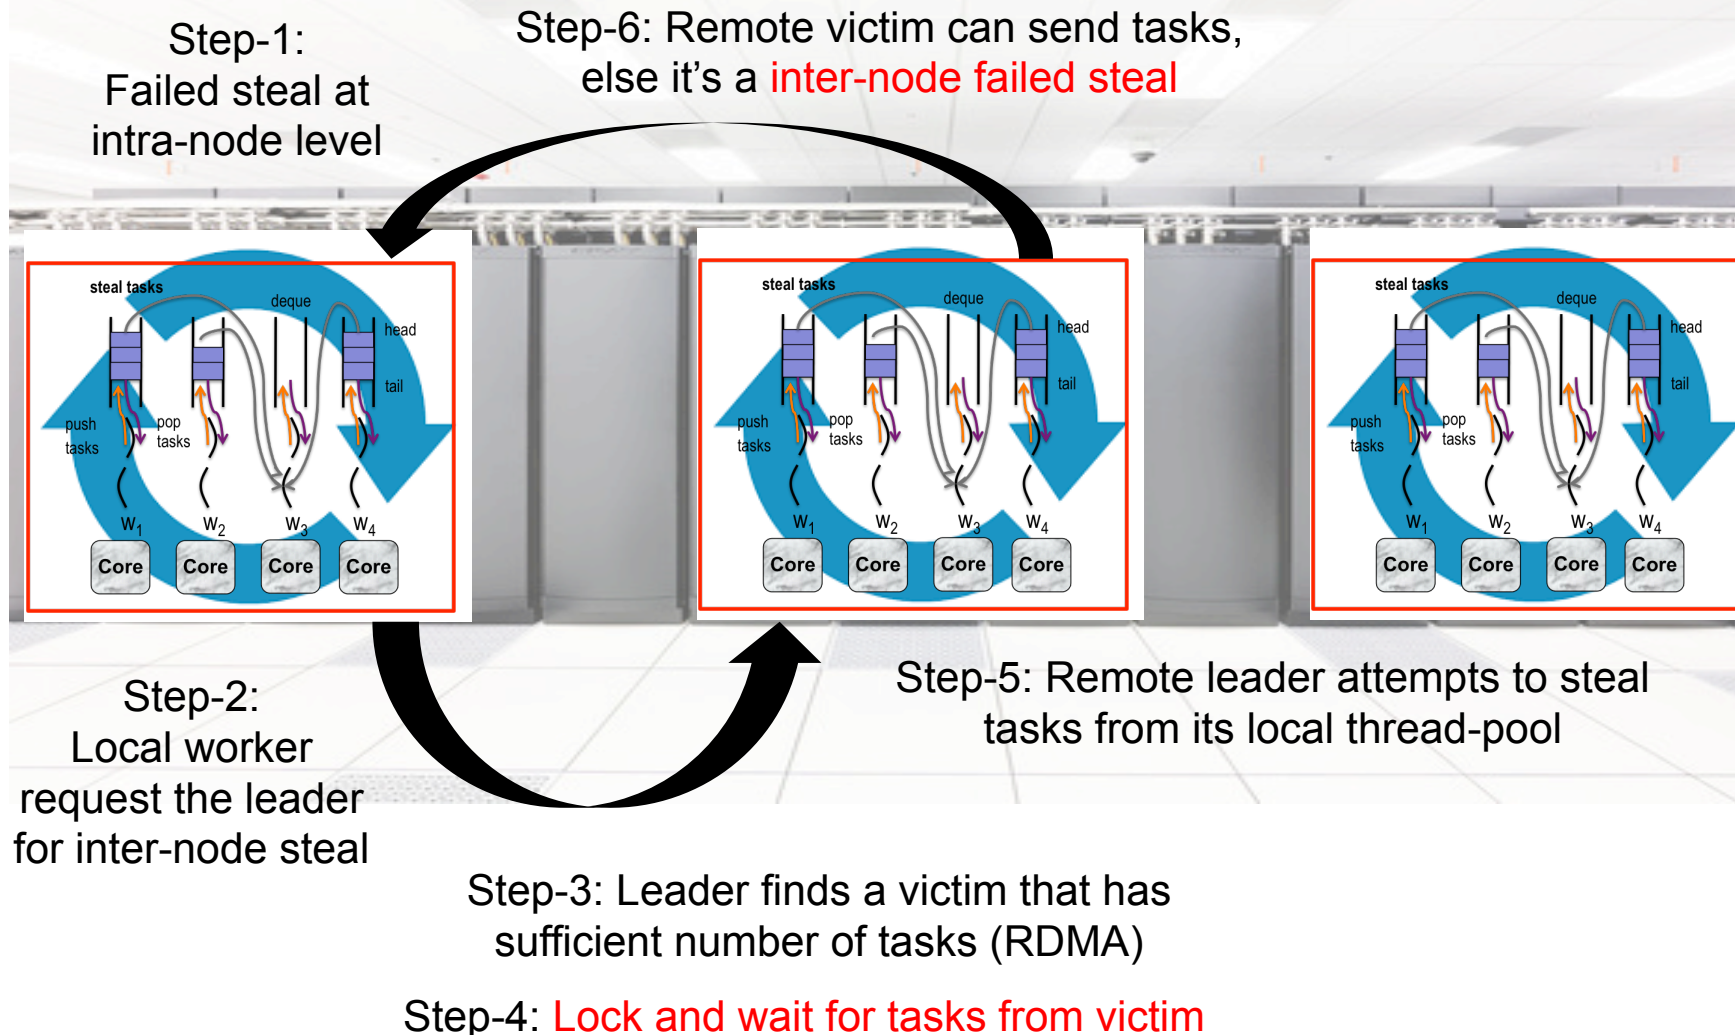
Distributed Work-Stealing Runtime

- Two different implementations in HabaneroUPC++
- BaselineWS
 - Uses prior work
- SuccessOnlyWS
 - Extends BaselineWS by using a novel victim selection policy that completely removes all inter-node failed steals




Implementation

BaselineWS in HabaneroUPC++



Implementation

SuccessOnlyWS in HabaneroUPC++



Methodology

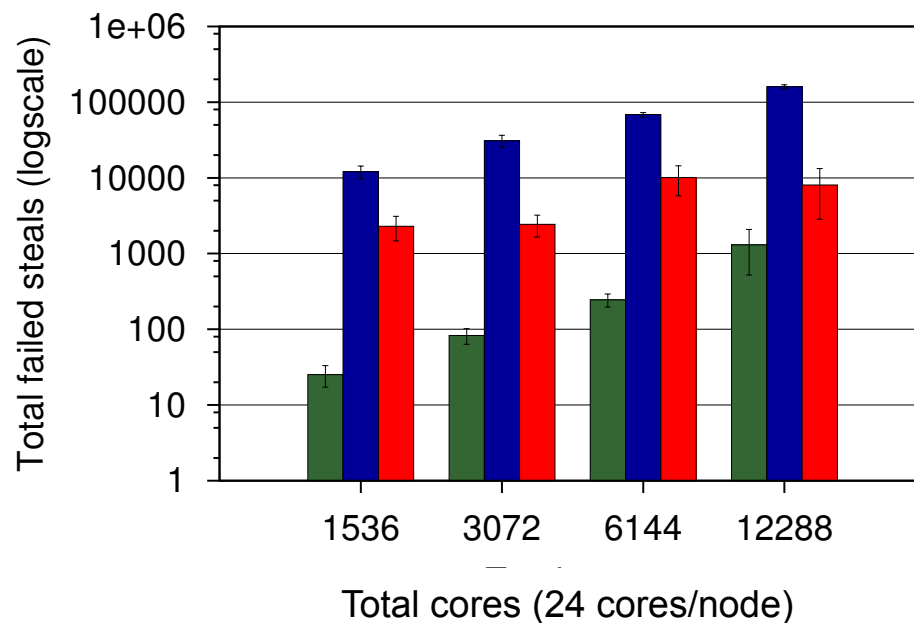
- Benchmarks
 - Two UTS trees T1WL and T3WL
 - NQueens
- Computing infrastructure
 - Edison supercomputer at NERSC
 - 2x12 cores per node



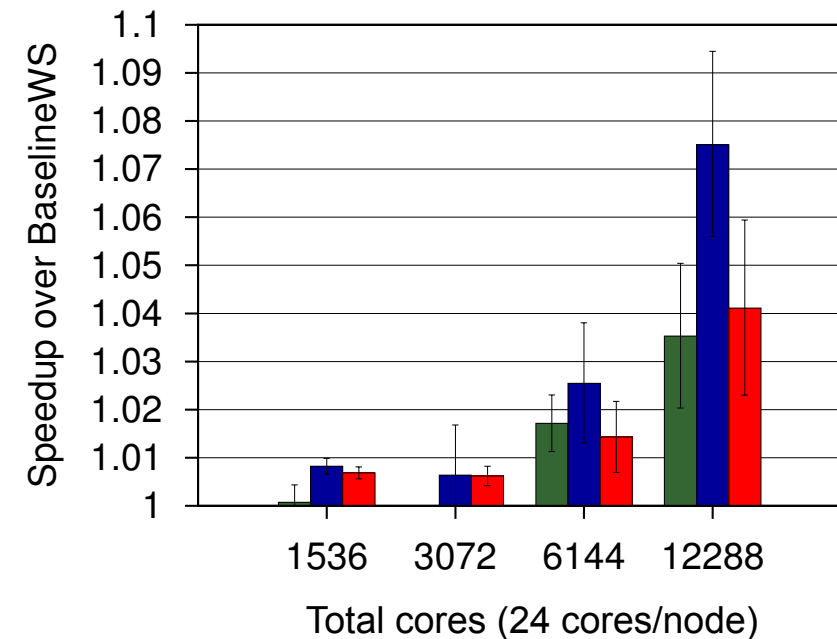
Experimental Evaluation

Results

Higher inter-node failed steals in BaselineWS \Rightarrow Better performance in SuccessOnlyWS



(a) Total inter-node failed steals in BaselineWS



(b) SuccessOnlyWS speedup over BaselineWS

T1WL

T3WL

NQueens

Note: More results are available in the paper



Summary and Conclusion

- Inter-node steals are costlier than intra-node steals
- Failed inter-node steals could hamper performance
- C++11 lambda function based API to in HabaneroUPC++ to express complex irregular computation that can participate in distributed work-stealing
- A novel implementation of distributed work-stealing runtime in HabaneroUPC++ PGAS library that completely removes all inter-node failed steals
- Our novel runtime delivers performance benefits up to 7%



Backup Slides



Existing Techniques for Inter-node Stealing

- Thread pool based hybrid runtimes [Lifflander et. al., HPDC'12, Paudel et. al., ICPP'13]
- Communication worker maintain ready queue of tasks even before a remote request arrives [Paudel et. al., ICPP'13]
- Load-aware steal attempts to *reduce* chances of failure [Dinan et. al., ICPP'08]
- First try random victims and on failing contact set of victims (lifelines) that promises to send tasks whenever they have it ready [Saraswat et. al., PPOPP'11]



Inter-node Steal Request from Thief

```
1 procedure Steal_AsyncAny
2     while (global termination is not detected)
3         V = get a random remote rank
4         if (V has declared task availability in PGAS space) // RDMA
5             if (I did not try to steal from V)
6                 queue my rank at V
7             if TryLock (V) is success
8                 save my rank at V
9                 wait until V send tasks or decline
10                Unlock (V)
11        break from while loop if I just received asyncAny tasks
12    if I receive asyncAny tasks from any victim
13        forget that I contacted this victim
14    reset my task receiving status
```

BaselineWS Runtime

SuccessOnlyWS Runtime



Inter-node Task Transfer from Victim

```
1 procedure Send_AsyncAny
2   while (there are pending inter-node steal requests)
3     T = get rank of one of the queued remote thief
4     steal tasks from my local workers and send to T
5     forget that T contacted me
6     break out of the while loop if local steal failed
7
8   T = get rank of the only waiting remote thief
9   steal tasks from local workers and send to T
10  declare that now I don't have any waiting remote thief
11  publish in PGAS space asyncAny count at my place
```

BaselineWS Runtime

SuccessOnlyWS Runtime

